

```
1 from google.colab import drive  
2 drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call dr

```
1 import pandas as pd  
2 import numpy as np  
3
```

```
1 df= pd.read_csv("/content/drive/MyDrive/Lab_performance/Breast Cancer Wiscc
```

```
1 #handle missing value  
2  
3 missing_counts = df.isna().sum()  
4 print(missing_counts[missing_counts > 0])  
5
```

Unnamed: 32 569
dtype: int64

```
1 #Duplicate rows remove  
2  
3 duplicate_count = df.duplicated().sum()  
4 print("Number of duplicate rows found:", duplicate_count)  
5 df_cleaned = df.drop_duplicates(keep='first')  
6 df_cleaned.reset_index(drop=True, inplace=True)  
7  
8 print("New shape after removing duplicates:", df_cleaned.shape)  
9  
10
```

Number of duplicate rows found: 0
New shape after removing duplicates: (569, 33)

```
1 # Correct data type  
2  
3 df['diagnosis'] = df.get('diagnosis').astype('category')  
4 print(df['diagnosis'].value_counts())  
5
```

diagnosis
B 357
M 212
Name: count, dtype: int64

```
1 # Identify categorical & numeric columns
2
3 cat_cols = df.select_dtypes(include=['object', 'category']).columns
4 num_cols = df.select_dtypes(include='number').columns
5
6 num_cols = num_cols.drop('id') if 'id' in num_cols else num_cols
7
8 print("Categorical columns:", list(cat_cols))
9 print("Total numeric columns:", len(num_cols))
10
```

```
Categorical columns: ['diagnosis']
Total numeric columns: 31
```

```
1 # Label Encoding for binary categorical features
2
3 from sklearn.preprocessing import LabelEncoder
4
5 if 'diagnosis' in df and df['diagnosis'].nunique() == 2:
6     df['diagnosis_le'] = LabelEncoder().fit_transform(df['diagnosis'])
7     print(df['diagnosis'].unique())
8
```

```
['M', 'B']
Categories (2, object): ['B', 'M']
```

```
1 #One-Hot Encoding for multi-class categorical features
2
3 from sklearn.preprocessing import OneHotEncoder
4 import pandas as pd
5
6 multi_class_cols = ['area_category', 'texture_group']
7 multi_class_cols = [col for col in multi_class_cols if col in df.columns]
8
9 if multi_class_cols:
10    print("Applying One-Hot Encoding on:", multi_class_cols)
11
12    ohe = OneHotEncoder(sparse=False, drop='first')
13    encoded = ohe.fit_transform(df[multi_class_cols])
14
15    encoded_df = pd.DataFrame(
16        encoded,
17        columns=ohe.get_feature_names_out(multi_class_cols),
18        index=df.index
19    )
20
21    df = df.drop(columns=multi_class_cols)
22    df = pd.concat([df, encoded_df], axis=1)
23
24 else:
25    print("⚠ No multi-class categorical columns found for One-Hot Encoding")
```

```
26  
27 df.info()  
28
```

⚠ No multi-class categorical columns found for One-Hot Encoding.

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 569 entries, 0 to 568

Data columns (total 34 columns):

#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	category
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64
16	smoothness_se	569 non-null	float64
17	compactness_se	569 non-null	float64
18	concavity_se	569 non-null	float64
19	concave points_se	569 non-null	float64
20	symmetry_se	569 non-null	float64
21	fractal_dimension_se	569 non-null	float64
22	radius_worst	569 non-null	float64
23	texture_worst	569 non-null	float64
24	perimeter_worst	569 non-null	float64
25	area_worst	569 non-null	float64
26	smoothness_worst	569 non-null	float64
27	compactness_worst	569 non-null	float64
28	concavity_worst	569 non-null	float64
29	concave points_worst	569 non-null	float64
30	symmetry_worst	569 non-null	float64
31	fractal_dimension_worst	569 non-null	float64
32	Unnamed: 32	0 non-null	float64
33	diagnosis_le	569 non-null	int64

dtypes: category(1), float64(31), int64(2)

memory usage: 147.5 KB

```
1 # Feature Scaling (StandardScaler)  
2  
3 from sklearn.preprocessing import StandardScaler  
4  
5 remove_cols = [ "diagnosis_le", "Unnamed: 32" ]  
6  
7 numeric_cols = [col for col in df.columns if col not in remove_cols]
```

```
8
9 numeric_cols = df[numeric_cols].select_dtypes(include='number').columns.to
10
11 print("Numeric Columns to Scale:", numeric_cols)
12
13 scaler = StandardScaler()
14 df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
15
16 print("Scaling Completed!")
17
```

```
Numeric Columns to Scale: ['id', 'radius_mean', 'texture_mean', 'perimeter_mean'
Scaling Completed!
```

```
1 #Train-Test Split
2
3 import pandas as pd
4 from sklearn.model_selection import train_test_split
5 from sklearn.preprocessing import StandardScaler, LabelEncoder
6
7 df= pd.read_csv("/content/drive/MyDrive/Lab_performance/Breast Cancer Wisc
8 df = df.loc[:, ~df.columns.str.contains("Unnamed")]
9
10 le = LabelEncoder()
11 df['diagnosis_le'] = le.fit_transform(df['diagnosis'])
12 print("Label Classes:", le.classes_)      # ['B' 'M']
13
14
15 numeric_cols = [c for c in df.select_dtypes(include='number').columns
16                  if c not in ['id', 'diagnosis_le']]
17 print("Numeric Columns to Scale:", numeric_cols)
18
19 scaler = StandardScaler()
20 df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
21 print("Scaling Completed!\n")
22
23 X = df.drop(columns=['id', 'diagnosis', 'diagnosis_le'])
24 y = df['diagnosis_le']
25
26 X_train, X_test, y_train, y_test = train_test_split(
27     X, y, test_size=0.2, stratify=y, random_state=42
28 )
29
30 print("Train shape:", X_train.shape, y_train.shape)
31 print("Test shape:", X_test.shape, y_test.shape)
32
33 print("\nTrain class distribution:\n",
34       y_train.value_counts(normalize=True))
35 print("\nTest class distribution:\n",
```

```
36     y_test.value_counts(normalize=True))  
37  
  
Label Classes: ['B' 'M']  
Numeric Columns to Scale: ['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean', 'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se', 'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se', 'fractal_dimension_se', 'radius_worst', 'texture_worst', 'perimeter_worst', 'area_worst', 'smoothness_worst', 'compactness_worst', 'concavity_worst', 'concave points_worst', 'symmetry_worst', 'fractal_dimension_worst']  
Scaling Completed!  
  
Train shape: (455, 30) (455,)  
Test shape: (114, 30) (114,)  
  
Train class distribution:  
diagnosis_le  
0    0.626374  
1    0.373626  
Name: proportion, dtype: float64  
  
Test class distribution:  
diagnosis_le  
0    0.631579  
1    0.368421  
Name: proportion, dtype: float64
```

```
1 #TRAIN FIVE ML MODELS  
2  
3 import pandas as pd  
4 from sklearn.model_selection import train_test_split  
5 from sklearn.preprocessing import StandardScaler  
6 from sklearn.metrics import accuracy_score  
7 from sklearn.linear_model import LogisticRegression  
8 from sklearn.svm import SVC  
9 from sklearn.neighbors import KNeighborsClassifier  
10 from sklearn.ensemble import RandomForestClassifier  
11 from sklearn.tree import DecisionTreeClassifier  
12 df= pd.read_csv("/content/drive/MyDrive/Lab_performance/Breast Cancer Wisconsin (Diagnostic).csv")  
13 df = df.loc[:, ~df.columns.str.contains('^\u0334nnamed')]  
14 df = df.fillna(df.median(numeric_only=True))  
15 X = df.drop("diagnosis", axis=1)  
16 y = df["diagnosis"]  
17 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
18 scaler = StandardScaler()  
19 X_train = scaler.fit_transform(X_train)  
20 X_test = scaler.transform(X_test)  
21 models = {  
22 "Logistic Regression": LogisticRegression(max_iter=500),  
23 "SVM": SVC(),  
24 "KNN": KNeighborsClassifier(),  
25 "Random Forest": RandomForestClassifier(),  
26 "Decision Tree": DecisionTreeClassifier()  
27 }  
28 for name, model in models.items():  
29     model.fit(X_train, y_train)  
30     pred = model.predict(X_test)
```

```
31     acc = accuracy_score(y_test, pred)
32     print(f"\n{name} Accuracy: {acc:.4f}")
```

```
Logistic Regression Accuracy: 0.9737
SVM Accuracy: 0.9825
KNN Accuracy: 0.9474
Random Forest Accuracy: 0.9649
Decision Tree Accuracy: 0.9123
```

```
1 #Confusion Matrix + Classification Report (Precision, Recall, F1-score)
2
3 from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
4 import seaborn as sns
5 import matplotlib.pyplot as plt
6
7 def evaluate_model(model, X_test, y_test, model_name):
8     y_pred = model.predict(X_test)
9
10    print("\n====")
11    print(f"Model: {model_name}")
12    print("====")
13
14    # Accuracy
15    acc = accuracy_score(y_test, y_pred)
16    print("Accuracy:", acc)
17
18    # Classification Report
19    print("\nClassification Report:")
20    print(classification_report(y_test, y_pred))
21
22    # Confusion Matrix
23    cm = confusion_matrix(y_test, y_pred)
24    plt.figure(figsize=(5,4))
25    sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
26    plt.title(f"\n{name} - Confusion Matrix")
27    plt.xlabel("Predicted")
28    plt.ylabel("Actual")
29    plt.show()
30
```


```
1 for model_name, model in models.items():
2     evaluate_model(model, X_test, y_test, model_name)
3
```



```
=====
Model: Logistic Regression
=====
Accuracy: 0.9736842105263158
```

Classification Report:

	precision	recall	f1-score	support
B	0.97	0.99	0.98	71
M	0.98	0.95	0.96	43
accuracy			0.97	114
macro avg	0.97	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114




```
=====
Model: SVM
=====
Accuracy: 0.9824561403508771
```

Classification Report:

	precision	recall	f1-score	support
B	0.97	1.00	0.99	71
M	1.00	0.95	0.98	43
accuracy			0.98	114
macro avg	0.99	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

SVM - Confusion Matrix



=====

Model: KNN


=====

Accuracy: 0.9473684210526315

Classification Report:

	precision	recall	f1-score	support
B	0.96	0.96	0.96	71
M	0.93	0.93	0.93	43
accuracy			0.95	114
macro avg	0.94	0.94	0.94	114
weighted avg	0.95	0.95	0.95	114

KNN - Confusion Matrix





```
=====
```

Model: Random Forest


```
=====
```

Accuracy: 0.9649122807017544

Classification Report:

	precision	recall	f1-score	support
B	0.96	0.99	0.97	71
M	0.98	0.93	0.95	43
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

Random Forest - Confusion Matrix



```
=====
```

Model: Decision Tree

```
=====
```

Accuracy: 0.9122807017543859

Classification Report:

	precision	recall	f1-score	support
B	0.94	0.92	0.93	71
M	0.87	0.91	0.89	43
accuracy			0.91	114
macro avg	0.90	0.91	0.91	114
weighted avg	0.91	0.91	0.91	114

```
1 #ROC Curve
2
3 from sklearn.metrics import roc_curve, auc
4
5 def plot_roc_curve(model, X_test, y_test, model_name):
6     # Probability predictions (required for ROC)
7     y_prob = model.predict_proba(X_test)[:, 1]
8
9     fpr, tpr, thresholds = roc_curve(y_test, y_prob)
10    roc_auc = auc(fpr, tpr)
11
12    plt.figure(figsize=(6,5))
13    plt.plot(fpr, tpr, label=f"{model_name} (AUC = {roc_auc:.3f})")
14    plt.plot([0,1], [0,1], linestyle="--")
15    plt.xlabel("False Positive Rate")
16    plt.ylabel("True Positive Rate")
17    plt.title(f"ROC Curve - {model_name}")
18    plt.legend()
19    plt.show()
```


```
1 from sklearn.preprocessing import LabelEncoder
2
3 le = LabelEncoder()
4 df['diagnosis'] = le.fit_transform(df['diagnosis'])
5
6 # Check mapping:
7 print(dict(zip(le.classes_, le.transform(le.classes_))))
8

{'B': np.int64(0), 'M': np.int64(1)}
```

```
1 X = df.drop("diagnosis", axis=1)
2 y = df["diagnosis"]
3
4 from sklearn.model_selection import train_test_split
5 X_train, X_test, y_train, y_test = train_test_split(
6     X, y, test_size=0.2, random_state=42
7 )
```


```
1 plot_roc_curve(models["Logistic Regression"], X_test, y_test, "Logistic Reg
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/utils/validation.py:2732: UserWarning:  
  warnings.warn(
```




```
1 #Loss vs Validation Loss Curve  
2 from sklearn.model_selection import learning_curve  
3 import numpy as np  
4  
5 def plot_learning_curve(model, X, y, model_name):  
6     train_sizes, train_scores, test_scores = learning_curve(  
7         model, X, y, cv=5, scoring='accuracy',  
8         train_sizes=np.linspace(0.1, 1.0, 10), n_jobs=-1  
9     )  
10    train_mean = train_scores.mean(axis=1)  
11    test_mean = test_scores.mean(axis=1)  
12  
13    plt.figure(figsize=(6,5))  
14    plt.plot(train_sizes, train_mean, label="Training Accuracy")  
15    plt.plot(train_sizes, test_mean, label="Validation Accuracy")  
16    plt.title(f"Learning Curve - {model_name}")  
17    plt.xlabel("Training Set Size")  
18    plt.ylabel("Accuracy")  
19    plt.legend()  
20    plt.show()
```

```
1 plot_learning_curve(models["Random Forest"], X, y, "Random Forest")
```



```
1 plot_learning_curve(models["Logistic Regression"], X, y, "Logistic Regressi
```



```
1 plot_learning_curve(models["Decision Tree"], X, y, "Decision Tree")
```

Learning Curve - Decision Tree

```
1 plot_learning_curve(models["SVM"], X, y, "SVM")
```

Learning Curve - SVM

The plot shows two curves: a blue line for Training Accuracy and an orange line for Validation Accuracy. Both curves start at approximately 0.9 accuracy for 10 samples and decrease as the number of samples increases to 100. The validation accuracy curve stays slightly above the training accuracy curve throughout.

Number of Samples	Training Accuracy	Validation Accuracy
10	~0.90	~0.90
20	~0.88	~0.89
30	~0.85	~0.86
40	~0.82	~0.83
50	~0.78	~0.80
60	~0.75	~0.78
70	~0.72	~0.75
80	~0.68	~0.72
90	~0.65	~0.69
100	~0.62	~0.66